

# Performance Analysis of RocksDB and its Integration into Open EdX

FUNDAMENTAL RESEARCH GROUP, IIT BOMBAY



# Our Team

---

- **Principle Investigator** : Prof Deepak B Phatak
- **Project in-Charge** : Mr. Nagesh Karmali & Miss Firuza Aibara
- **Mentor** : Nithin S
- Project Interns: Abhishek Chattopadhyay  
Ekta Agrawal  
Simran Goyal

# Primary Objective

---

- Performance improvement of OpenEdX using Facebook's RocksDB engine.
- Performance Comparison between InnoDB and RocksDB and feasibility study.

# Software Prerequisites

---

- OpenEdX
- JMeter
- MariaDB with RocksDB Plugin

# JMeter

---

Apache JMeter is an open source software, Java application designed to load test functional behavior and measure performance.

# OpenEdX

OpenEdX is an open source MOOC Platform for providing online courses.

It consists of 2 components:

- LMS
- CMS

MySQL is in the LMS to store User Content

# MariaDB

An Enhanced Drop-In replacement for MySQL, which is faster and has a richer eco-system of storage engines and plugins as compared to MySQL

# RocksDB

- RocksDB is an embedded, high performance, persistent key-value storage engine developed at Facebook Inc.
- RocksDB uses Log-Structured Merge trees (LSM) to obtain significant space efficiency and better write throughput.
- It is used by major organisations like Facebook, Yahoo!, LinkedIn and Netflix




RocksDB

# Compaction in Log-Structured Merge Tree

---

- The LSM Tree is an N-level tree with levels  $C_0$ - $C_N$ , where  $C_0$  is always present in Memory and  $C_1$  to  $C_n$  in the disk.
- Insertion request to database enters into the  $C_0$  level and is subsequently transferred to  $C_1$  and the corresponding lower levels.
- This transfer happens through the process of a “Rolling Merge”, similar to Merge Sort.
- Each merge reads a leaf node of the  $(l+1)^{\text{th}}$  level buffered in the block, merges it with the entries of the current level ( $l$ ) and creates a newly merged node of the  $(l+1)^{\text{th}}$  level, which are returned to new blocks so that the old block is not overwritten. This is the compaction process in LSM.

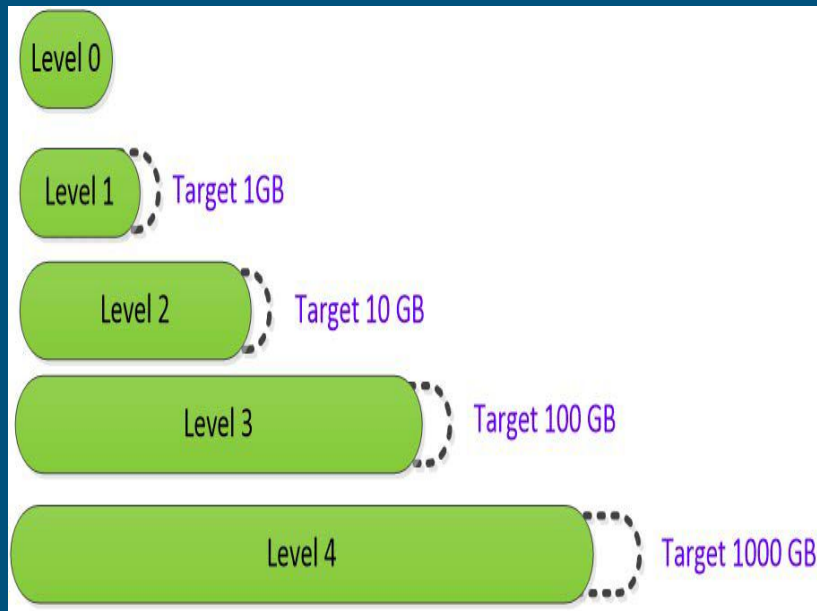
# Compaction Procedure





# LSM Compaction

- For each level, data is sorted by key
- LSM Amplification depends on the stale data yet to be collected.
- In the worst case, LSM-tree space amplification for data is 1.11, considering that size of all levels until the last level amounts to only 11.11% of the last level. (Default Case)
- The level size multiplier is a tunable parameter which can be used to moderate the space, read and write amplification.



# Machine Specifications

---

- Remote testing of two Virtual Machines (VM1 and VM2) using JMeter.
- Configuration of the VM's:
  - 4 cores x 4 sockets
  - 2GB RAM
  - 230 GB Hard Drive
  - Ubuntu Server 16.04 LTS

# Test Parameters

---

Property	VM1	VM2
innodb_buffer_pool_size	256	512
innodb_log_file_size	128	256
innodb_buffer_pool_instances	4	8
max_connections	512	1024
innodb_spin_wait_delay	0	0

# Test Procedure

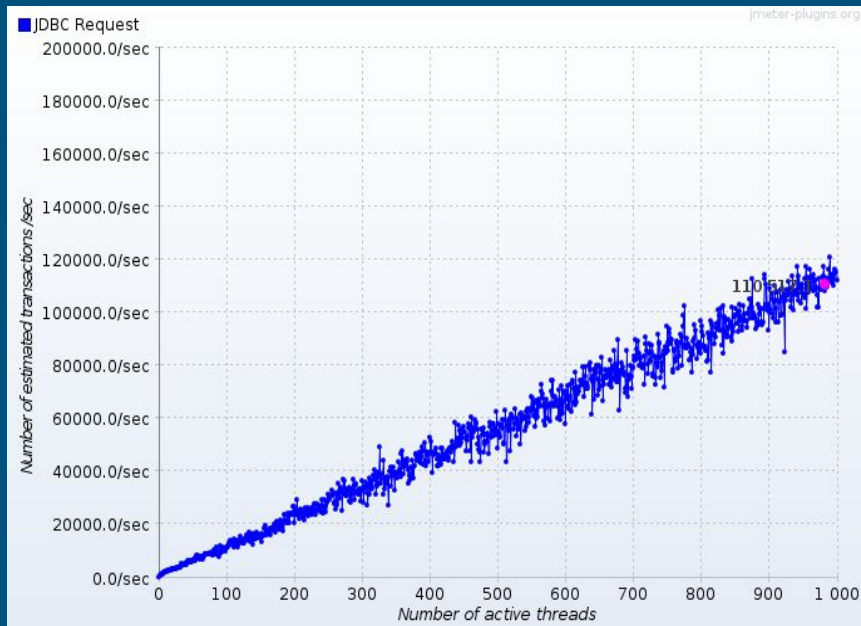
---

Set Up and Tune JMeter -> Make Changes to '*my.cnf*' file in VM -> Add required test plan -> Set up JDBC Configuration -> Select the type of JDBC Request -> Save and Run Test Plan in Command Line Mode


Open JMeter (GUI) -> Add Listeners to Workbench -> Analyse results

# Graph Analysis for Insert Query

Number of Threads : 1000 Ramp-Up Time : 100



InnoDB



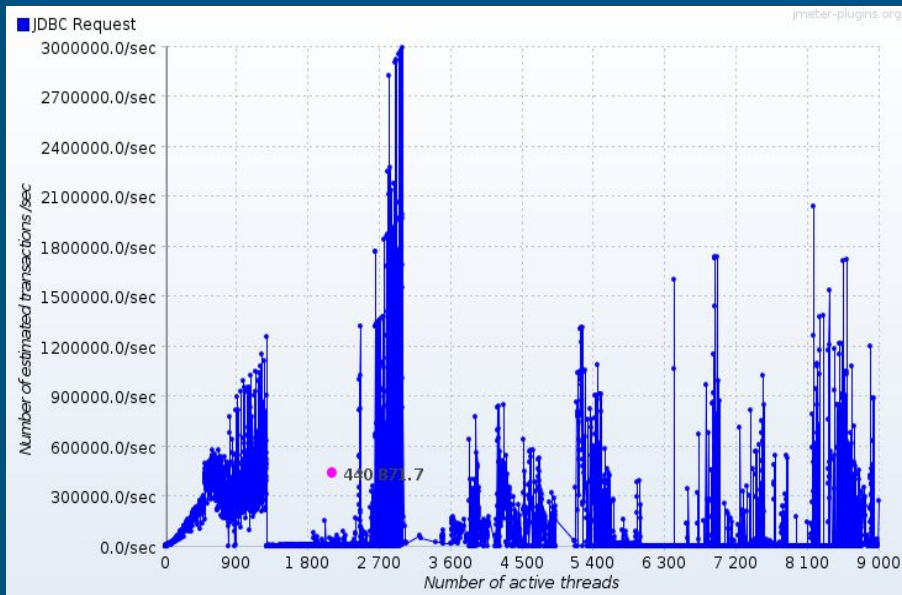
RocksDB

# Insert Query Specifications


Basis	InnoDB	RocksDB
Bytes Throughput	24500 Bytes/Sec	28000 Bytes/Sec
Response Latencies	500 ms	390 ms
Transaction Throughput	110517 Transactions/thread	175525 Transactions/thread
Transactions per Second	2150 Transactions/sec	2700 Transactions/sec

# Graph Analysis For Select Query

Number of Threads : 9000 Ramp-Up Time : 300



InnoDB



RocksDB

# Select Query Specifications

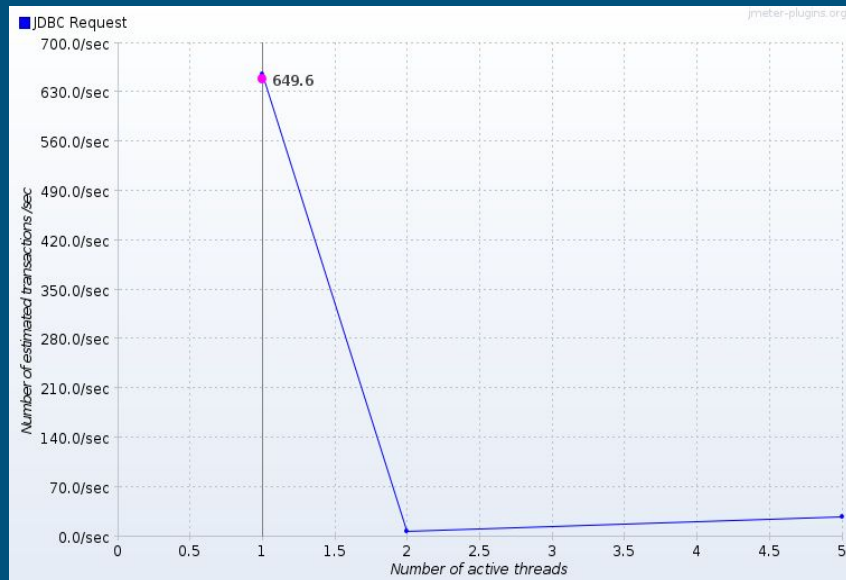
Basis	InnoDB	RocksDB
Bytes Throughput	27000	27000
Response Latencies	6-8	7-8
Transaction Throughput	440871 Transactions/ threads	436169 Transactions/ threads
Transactions per Second	4700 Transactions/sec	4500 Transactions/sec




# Graph Analysis For Update Query

Number of Threads : 500

Ramp-Up Time : 50



**InnoDB**



**RocksDB**

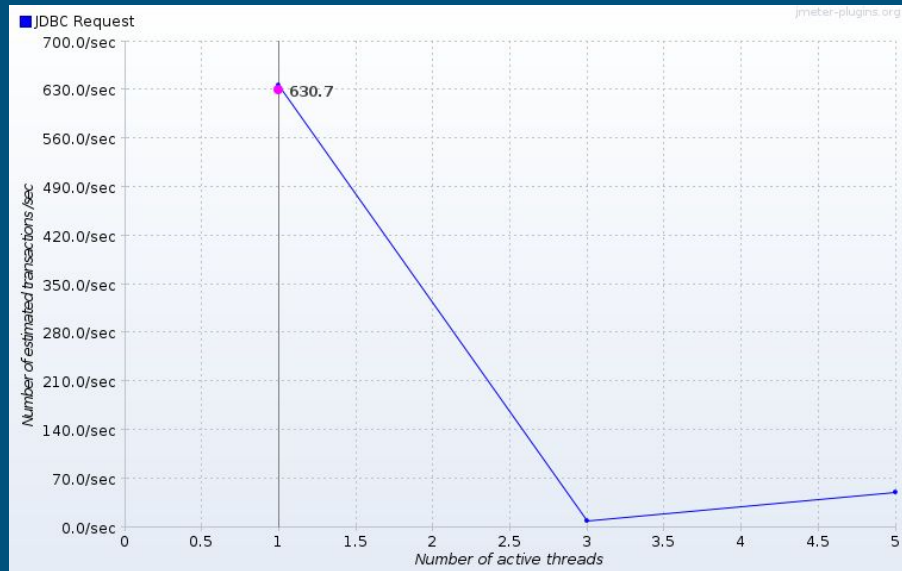
# Update Query Specifications

Basis	InnoDB	RocksDB
Bytes Throughput	110 Bytes/sec	110 Bytes/sec
Response Latencies	3 ms	3-4 ms
Transaction Throughput	649 Transactions/thread	614 Transactions/thread
Transactions per Second	10 Transactions/sec	10 Transactions/sec


# Graph Analysis For Delete Query

Number of Threads : 500

Ramp-Up Time : 50



InnoDB



RocksDB

# Delete Query Specifications

Basis	InnoDB	RocksDB
Bytes Throughput	110 Bytes/sec	110 Bytes/sec
Response Latencies	3 ms	2 ms
Transaction Throughput	630 Transactions/ thread	690 Transactions/ thread
Transactions per Second	10 Transactions/sec	10 Transactions/sec

# Relative ORDER BY clause time frame

---

Database Name	Time Taken	Throughput
University	50 sec	9.9 Threads/sec
Universityrocks	10 min 33 sec	0.8 Threads/sec

Number of Threads : 500    Ramp-Up Time : 50

# Relative Database Sizes

Database Name	Database Size (MB)
University	4540.54687500
Universityrocks	1361.6663660

VM1

Database Name	Database Size (MB)
University	4611.51562500
Universityrocks	1254.96938515

VM2

# Test Analysis

---

- High Compression Efficiency (~66% Smaller Database)
- Improved Performance in INSERT and DELETE queries by RocksDB
- Comparable Performance in simple SELECT and UPDATE queries
- Lags behind in ORDER BY due to underlying data structure.
- Lack of FOREIGN KEY support in RocksDB

# Final Results

---

- There is noticeable improvement in performance for RocksDB over InnoDB with regards to Storage Space efficiency and higher average Transaction Throughput for queries which relate to write operations
- Integration of RocksDB with OpenEdX is currently not feasible, since it does not support FOREIGN KEY constraint.



# Future Works

---

- Implementation of FOREIGN KEY in RocksDB
- RocksDB compilation with MySQL
- Implementation of SAVEPOINT in MySQL
- Improving the performance of ORDER BY Queries
- Integration of OpenEdX with RocksDB

- 
- For detailed analysis and graphs, kindly refer to our GitHub repository:

<https://github.com/fresearchgroup/performance-improvement-of-openedx>

<https://github.com/fresearchgroup/performance-improvement-of-openedx/wiki>



Thank you!

## Absence of bottleneck issues in MongoDB

